Enterprise RAG Knowledge Assistant

Retrieval-Augmented Generation (RAG) combines document retrieval with a language model to answer questions using relevant context.

Our system stores document chunks and their 768-dimensional embeddings in PostgreSQL with the pgvector extension.

Ollama runs the language model and embedding model locally, without requiring a paid external API.